

Nilla AI Agent Architecture

This document describes the AI agent system powering Nilla's intelligent features. These are task-specific agents designed to help open source contributors succeed, not general-purpose chatbots.

Why Agents, Not a Chatbot

Nilla uses **specialized agents** rather than a conversational chatbot for several important reasons:

1. Task-Oriented Design

Each agent is purpose-built for a specific job. They don't engage in open-ended conversation, they receive structured inputs, make decisions, and return structured outputs. This makes them:

- **Predictable:** Same inputs produce consistent, reliable outputs
- **Testable:** Structured I/O enables automated testing and validation
- **Focused:** No risk of going off-topic or hallucinating unrelated content

2. Structured Input/Output Contracts

Every agent has Zod-validated schemas for both input and output. This ensures:

- Type safety throughout the application
- Clear API contracts between frontend and AI
- Graceful fallbacks when LLM output is malformed

3. Domain-Specific Prompting

Each agent has carefully crafted system prompts with:

- Specific personas (e.g., "supportive commitment coach")
- Clear rules and constraints
- Output format specifications
- Context-appropriate behavior based on user state

4. Deterministic Logic + LLM Reasoning

Agents combine deterministic calculations (time remaining, progress percentage, risk levels) with LLM reasoning (personalized messages, explanations). This hybrid approach provides:

- Reliable metrics the UI can depend on
- Human-like coaching and explanation quality
- Auditability of the decision-making process

Agent Overview

Agent	Purpose	RAG-Enabled
Issue Explainer	Breaks down GitHub issues for contributors at their skill level	Yes
Commitment Coach	Provides personalized guidance to keep contributors on track	No

Issue Recommender	Suggests the best-fit issue from a list of candidates	No
--------------------------	---	----

Issue Explainer Agent

Purpose: Help contributors understand GitHub issues by providing plain-English explanations, expected outcomes, and actionable guidance tailored to their experience level.

Inputs

```
interface IssueExplainerInput {
  issue: {
    title: string;           // Issue title
    body?: string;          // Issue description (truncated to 3000 chars)
    labels: string[];        // Labels like "bug", "good-first-issue"
    repository: string;     // "owner/repo" format
    url: string;            // Link to the issue
  };
  user: {
    experienceLevel: "beginner" | "intermediate" | "advanced";
  };
  repoContext?: string;      // RAG-retrieved documentation (optional)
}
```

Outputs

```
interface IssueExplainerOutput {
  summary: string;           // Plain-English explanation of the issue
  expectedOutcome: string;    // What "done" looks like
  repoGuidelines: string[];   // Repo-specific rules from documentation
  beginnerPitfalls: string[]; // Common mistakes to avoid
  suggestedApproach: string; // Step-by-step guidance
  keyTerms: Array<{
    term: string;
    definition: string;
  }>;
  confidenceNote: string;     // What to verify with maintainers
}
```

Decisions Made

1. **Explanation Depth:** Adapts language complexity based on experience level
 - Beginners: Simple language, defines jargon, assumes no Git/PR knowledge
 - Intermediate: Clear but not basic, focuses on project-specific conventions
 - Advanced: Concise, skips basics, highlights nuances
2. **Guideline Extraction:** Only includes repo guidelines that actually appear in the provided RAG context—never invents rules

3. **Pitfall Identification:** Proactively warns about common mistakes relevant to the issue type and user level

4. **Confidence Assessment:** Distinguishes between what the agent is confident about vs. what the contributor should verify

RAG Integration

The Issue Explainer is the **only agent that uses RAG**. When invoked:

1. The API fetches the tracked repository's ID
2. Calls `retrieveRepoContext()` with the issue title and body
3. Performs vector similarity search against ingested repo documentation
4. Injects matching chunks into the prompt as `repoContext`

If RAG context is available, the agent extracts specific contribution guidelines. If not, `repoGuidelines` is explicitly set to an empty array to prevent hallucination.

Commitment Coach Agent

Purpose: Provide personalized, actionable coaching to help contributors stay on track with their 7-day commitments.

Inputs

```
interface CommitmentCoachInput {  
  commitment: {  
    id: string;  
    issueTitle: string;  
    issueUrl: string;  
    repository: string;  
    createdAt: string;           // ISO timestamp  
    deadlineAt: string;         // ISO timestamp (7 days from creation)  
    currentMilestone: Milestone;  
    milestonesCompleted: Milestone[];  
    lastActivityAt?: string;  
  };  
  user: {  
    username: string;  
    totalCommitments?: number;  
    completedCommitments?: number;  
    timezone?: string;  
  };  
  currentTime?: string;          // For testing/reproducibility  
}  
  
type Milestone =  
  | "not_started"  
  | "read_issue"  
  | "ask_question"  
  | "work_on_solution"
```

```
| "open_pr"  
| "completed";
```

Outputs

```
interface CommitmentCoachOutput {  
    nextAction: {  
        action: string; // Specific next step to take  
        why: string; // Why this matters now  
        estimatedMinutes?: number; // Time estimate  
    };  
    nudge: {  
        message: string; // Encouraging coach message  
        tone: "encouraging" | "motivating" | "celebratory" | "urgent" | "supportive";  
    };  
    riskAssessment: {  
        level: "on_track" | "needs_attention" | "at_risk" | "critical";  
        reason: string;  
        daysRemaining: number;  
        hoursRemaining: number;  
    };  
    warning?: { // Only present if at_risk or critical  
        message: string;  
        suggestion: string;  
    };  
    progress: {  
        currentMilestone: Milestone;  
        milestonesRemaining: number;  
        percentComplete: number; // 0-100  
    };  
    meta: {  
        generatedAt: string;  
        commitmentId: string;  
    };  
}
```

Decisions Made

1. Risk Level Calculation (Deterministic):

- o `critical` : Overdue OR <24 hours remaining and not working on solution yet
- o `at_risk` : <2 days remaining and haven't started working
- o `needs_attention` : <4 days remaining and still in early stages
- o `on_track` : Everything else

2. Tone Selection: LLM chooses based on situation:

- o `encouraging` : Normal progress
- o `motivating` : Slow progress, needs a push
- o `celebratory` : Good progress, milestones achieved
- o `urgent` : Critical deadline situations

- o supportive : Struggling users who need empathy
3. **Next Action Specificity:** Provides concrete actions based on current milestone, not vague encouragement
4. **Warning Generation:** Only surfaces warnings when risk is elevated, avoiding alarm fatigue

Why No RAG

The Commitment Coach doesn't use RAG because:

- It focuses on time management and accountability, not technical details
- All necessary context (deadlines, milestones, progress) is passed directly
- Repo-specific guidance is already provided by the Issue Explainer

Issue Recommender Agent

Purpose: Analyze a list of candidate issues and recommend the best fit for a specific user based on their skills, interests, and experience.

Inputs

```
interface RecommendIssueInput {
  user: {
    id: string;
    username: string;
    skillLevel: "beginner" | "intermediate" | "advanced";
    preferredLanguages: string[];
    interests?: string[]; // e.g., ["testing", "docs", "frontend"]
    pastContributions?: number;
    availableHoursPerWeek?: number;
  };
  issues: Array<{ // 1-20 candidate issues
    id: string;
    title: string;
    body?: string;
    labels: string[];
    repository: string;
    language?: string; // Primary repo language
    openedAt?: string;
    commentCount?: number;
    url: string;
  }>;
}
```

Outputs

```
interface RecommendIssueOutput {
  recommendedIssue: {
    id: string;
    title: string;
```

```

    repository: string;
    url: string;
};

explanation: string; // Why this is the best fit
riskLevel: "low" | "medium" | "high";
riskFactors: string[]; // Specific concerns
alternativeIssues: Array<{
    id: string;
    title: string;
    reason: string;
}>;
rankedIssues: Array<{
    issueId: string;
    difficultyScore: number; // 1-10
    fitScore: number; // 1-10
    reasoning: string;
}>;
}

```

Decisions Made

1. Difficulty Assessment: Scores each issue 1-10 based on:

- Labels (good-first-issue = easier)
- Description complexity
- Comment volume (more comments may indicate complexity)
- Repository language match

2. User Fit Analysis: Scores each issue 1-10 based on:

- Skill level alignment
- Language proficiency match
- Interest alignment
- Time commitment vs. availability

3. Risk Evaluation:

- low : Clear scope, matches skills, maintainer responsive
- medium : Some ambiguity, may require learning
- high : Unclear requirements, skill gap, potential scope creep

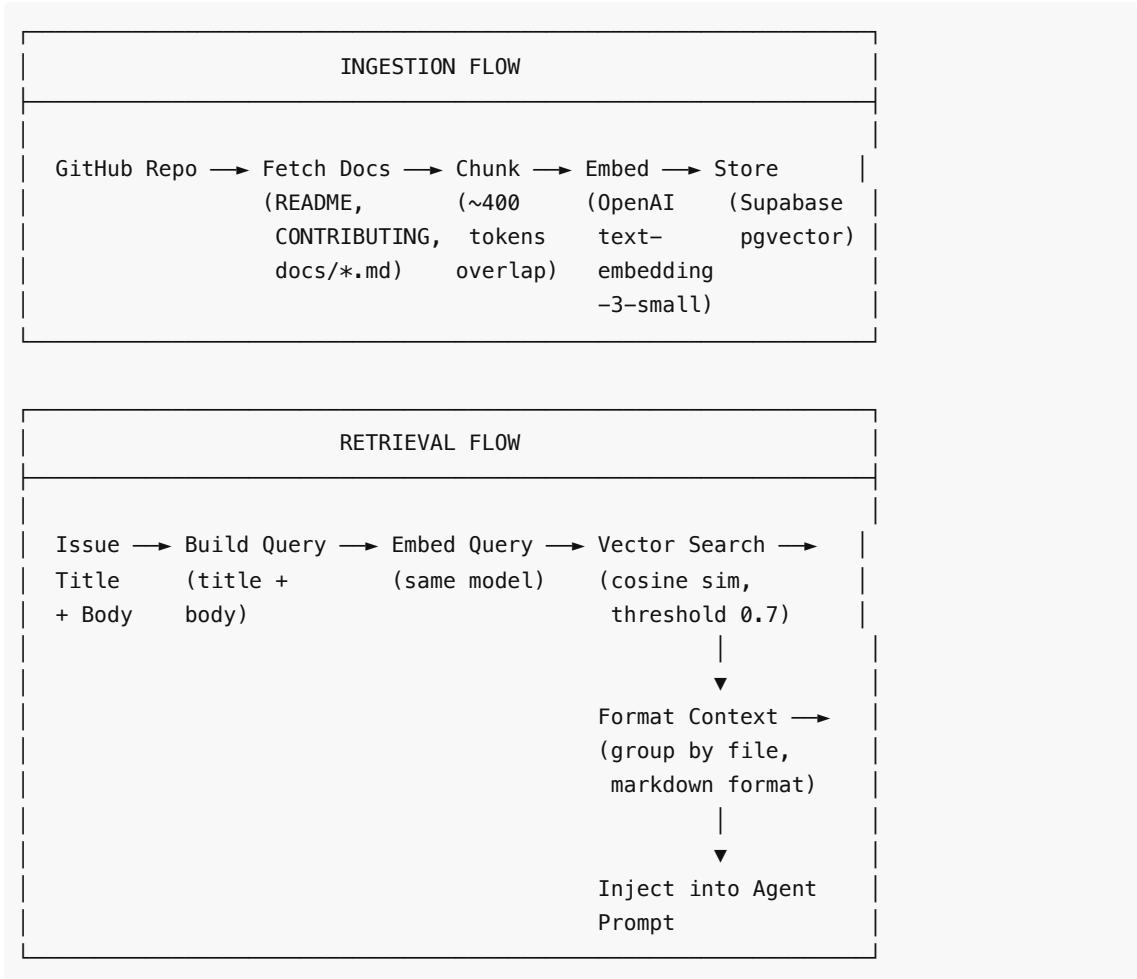
4. Ranking: All issues ranked by overall suitability, not just difficulty

5. Alternative Suggestions: Provides up to 2 alternatives with reasoning

RAG System Architecture

The RAG (Retrieval-Augmented Generation) system provides repository-specific context to the Issue Explainer agent.

Pipeline Overview



Ingestion Details

Triggered: When a repository is added/tracked

Documents Fetched:

- README.md
- CONTRIBUTING.md
- docs/*.md files

Chunking Strategy:

- Target size: 400 tokens (1600 characters)
- Overlap: 50 tokens between chunks
- Split on paragraph boundaries for coherence
- Each chunk retains file path and index metadata

Embedding Model: text-embedding-3-small (1536 dimensions)

Storage: Supabase `repo_document_chunks` table with pgvector extension

Retrieval Details

Query Construction: Issue title + body (truncated to 8000 chars)

Similarity Search:

- Cosine similarity via `match_repo_documents` RPC
- Default threshold: 0.7 (configurable)
- Default match count: 5 chunks

Context Formatting:

- Chunks grouped by source file
- Sorted by original position within each file
- Formatted as markdown for prompt injection

Observability with Opik

All AI agent calls are traced using Opik for observability:

```
const trackedAI = createTrackedAI("issue-explainer-completion");

const completion = await trackedAI.chat.completions.create({
  model: "gpt-4o",
  messages: [...],
});

await flushTraces();
```

This enables:

- Cost tracking per agent type
- Latency monitoring
- Prompt/response logging for debugging
- Quality analysis over time

File Structure

```
lib/ai/
├── openai.ts          # OpenAI client setup + Opik tracking
├── index.ts           # Type exports
└── agents/
    ├── index.ts        # Re-exports all agents
    ├── issue-explainer.ts # Issue Explainer agent
    ├── commitment-coach.ts # Commitment Coach agent
    └── recommend-issue.ts # Issue Recommender agent

lib/rag/
├── ingest.ts          # Document chunking + embedding pipeline
├── retrieve.ts         # Vector search + context formatting
└── fetch-repo-docs.ts # GitHub documentation fetching
```

Design Principles

1. **Fail Gracefully:** Every agent has fallback responses when LLM parsing fails
2. **Never Hallucinate Rules:** Agents only reference guidelines present in provided context
3. **Deterministic Where Possible:** Calculations (risk, progress, time) are done in code, not by the LLM
4. **Single Responsibility:** Each agent does one job well
5. **Structured Everything:** Zod schemas for inputs, outputs, and intermediate types
6. **Trace Everything:** All LLM calls are observable for debugging and improvement